World Models and Mesa-Optimisation: Mathematical Foundations for Agentic AI

Research Notes — ETH Zürich ML Lecture Series

Spring Semester 2026

Abstract

These notes develop two threads that converge on a central question for agentic AI systems: *how agents build internal representations of their environment, and what happens when the objectives pursued by those internal representations diverge from the objectives intended by their designers.* We first formalise **world models** in model-based reinforcement learning—covering the variational sequence model (RSSM), the Dreamer family (v1–v3, DayDreamer), and MuZero—with full derivations of the sequential ELBO, λ-return estimates, and value-equivalent modelling. We then pivot to the AI-safety landscape: **mesa-optimisation** (Hubinger et al., 2019), **reward hacking** (Skalse et al., 2022), and **goal misgeneralisation** (Langosco et al., 2022). We close by connecting these ideas to modern large language models viewed as implicit world models and simulators, articulate open problems, and provide an annotated bibliography of ≥ 15 key references.

Contents

1 Introduction **3**

2 World Models: Mathematical Foundations **3**
- 2.1 Setup and Notation . . . 3
- 2.2 Deriving the Sequential ELBO . . . 4
- 2.3 Information Bottleneck Interpretation . . . 4
- 2.4 Recurrent State-Space Model (RSSM) . . . 4
- 2.5 Historical Context . . . 5

3 The Dreamer Family **5**
- 3.1 Dreamer v1: Actor-Critic in Imagination . . . 5
  - 3.1.1 Imagination-Based Training . . . 5
  - 3.1.2 λ-Returns . . . 6
- 3.2 Dreamer v2: Discrete Representations and KL Balancing . . . 6
- 3.3 Dreamer v3: Symlog Predictions and Scale-Free Learning . . . 6
- 3.4 DayDreamer: Sim-to-Real via Learned World Models . . . 7

4 MuZero and Value-Equivalent Modelling **7**
- 4.1 Three Learned Functions . . . 7
- 4.2 Training Loss . . . 7
- 4.3 MCTS with a Learned Model . . . 8
- 4.4 Value-Equivalent Models . . . 8

5 Mesa-Optimisation **8**
5.1 Definitions (Hubinger et al., 2019) . . . 8
5.2 Why Mesa-Optimisers Arise . . . 9
5.3 Deceptive Alignment . . . 9
5.4 Taxonomy of Inner Misalignment . . . 9

6 Reward Hacking and Goodhart's Law **10**
6.1 Formal Setup (Skalse et al., 2022) . . . 10
6.2 Goodhart's Taxonomy: Formal Characterisation . . . 10
6.2.1 Regressional Goodhart . . . 10
6.2.2 Extremal Goodhart . . . 10
6.2.3 Causal Goodhart . . . 10
6.2.4 Adversarial Goodhart . . . 11
6.3 Implications for RLHF . . . 11

7 Goal Misgeneralisation **11**
7.1 Formal Setup (Langosco et al., 2022) . . . 11
7.2 The CoinRun Example . . . 11
7.3 Connection to Causal Inference . . . 12
7.4 Danger for Agentic Systems . . . 12

8 Connections to Modern Agentic AI **12**
8.1 LLMs as World Models . . . 12
8.2 "LLM as Simulator" (Janus, 2022) . . . 12
8.3 In-Context Learning as Mesa-Optimisation . . . 13
8.4 Scaling Risk . . . 13

9 Open Problems **13**

10 Annotated Bibliography **14**

1 Introduction

Model-based reinforcement learning (MBRL) promises sample efficiency by learning a *world model*—a compact, differentiable surrogate of the environment—and planning or training policies inside it. At the same time, as optimisation pressure increases, learned systems may develop internal structure that itself performs optimisation, potentially pursuing objectives misaligned with those of the outer training procedure.

This document is organised as follows.

1. Section 2: Mathematical foundations of world models.
2. Section 3: The Dreamer family of algorithms.
3. Section 4: MuZero and value-equivalent modelling.
4. Section 5: Mesa-optimisation and inner alignment.
5. Section 6: Reward hacking and Goodhart's taxonomy.
6. Section 7: Goal misgeneralisation.
7. Section 8: Connections to modern agentic AI.
8. Section 9: Open problems.
9. Section 10: Annotated bibliography.

2 World Models: Mathematical Foundations

2.1 Setup and Notation

Consider a partially observable Markov decision process (POMDP) $\mathcal{M} = (\mathcal{S}, \mathcal{A}, \mathcal{O}, T, O, R, \gamma)$ where:

- $\mathcal{S}$ is the (hidden) state space,
- $\mathcal{A}$ is the action space,
- $\mathcal{O}$ is the observation space,
- $T(s_{t+1} \mid s_t, a_t)$ is the transition model,
- $O(o_t \mid s_t)$ is the observation (emission) model,
- $R(r_t \mid s_t, a_t)$ is the reward model,
- $\gamma \in [0, 1)$ is the discount factor.

A **world model** learns parametric approximations $p_\theta(s_{t+1} \mid s_t, a_t)$, $p_\theta(o_t \mid s_t)$, and $p_\theta(r_t \mid s_t, a_t)$ from interaction data.

2.2 Deriving the Sequential ELBO

Given a trajectory $\tau = (o_1, a_1, r_1, \ldots, o_T, a_T, r_T)$, we wish to maximise the marginal log-likelihood:

$$\log p_\theta(\tau) = \log \int \prod_{t=1}^{T} p_\theta(o_t \mid s_t)\, p_\theta(r_t \mid s_t, a_t)\, p_\theta(s_t \mid s_{t-1}, a_{t-1})\, ds_{1:T}. \tag{1}$$

The integral is intractable, so we introduce a variational posterior $q_\phi(s_{1:T} \mid o_{1:T}, a_{1:T})$ that factorises as

$$q_\phi(s_{1:T} \mid o_{1:T}, a_{1:T}) = \prod_{t=1}^{T} q_\phi(s_t \mid s_{t-1}, a_{t-1}, o_t). \tag{2}$$

Theorem 2.1 (Sequential ELBO)**.** *Under the factorisations above,*

$$\log p_\theta(\tau) \;\geq\; \underbrace{\sum_{t=1}^{T} \mathbb{E}_{q_\phi}\Big[\log p_\theta(o_t \mid s_t) + \log p_\theta(r_t \mid s_t, a_t)\Big]}_{\textit{reconstruction}} \;-\; \underbrace{\sum_{t=1}^{T} \mathbb{E}_{q_\phi}\Big[\mathrm{KL}\big(q_\phi(s_t \mid s_{t-1}, a_{t-1}, o_t) \,\|\, p_\theta(s_t \mid s_{t-1}, a_{t-1})\big)\Big]}_{\textit{complexity}} \tag{3}$$

Proof. Starting from (1), apply Jensen's inequality via the variational distribution:

$$\begin{aligned}
\log p_\theta(\tau) &= \log \int q_\phi(s_{1:T} \mid o_{1:T}, a_{1:T})\, \frac{\prod_t p_\theta(o_t \mid s_t)\, p_\theta(r_t \mid s_t, a_t)\, p_\theta(s_t \mid s_{t-1}, a_{t-1})}{q_\phi(s_{1:T} \mid o_{1:T}, a_{1:T})}\, ds_{1:T} && (4)\\
&\geq \int q_\phi(s_{1:T} \mid o_{1:T}, a_{1:T})\, \log \frac{\prod_t p_\theta(o_t \mid s_t)\, p_\theta(r_t \mid s_t, a_t)\, p_\theta(s_t \mid s_{t-1}, a_{t-1})}{\prod_t q_\phi(s_t \mid s_{t-1}, a_{t-1}, o_t)}\, ds_{1:T} && (5)\\
&= \sum_{t=1}^{T} \mathbb{E}_{q_\phi}\Big[\log p_\theta(o_t \mid s_t) + \log p_\theta(r_t \mid s_t, a_t) - \log \frac{q_\phi(s_t \mid s_{t-1}, a_{t-1}, o_t)}{p_\theta(s_t \mid s_{t-1}, a_{t-1})}\Big]. && (6)
\end{aligned}$$

Recognising the last term as the KL divergence completes the proof. □

2.3 Information Bottleneck Interpretation

The KL term in (3) acts as an **information bottleneck**: it penalises the posterior q_ϕ for encoding more information from o_t than is strictly necessary for prediction. This connects world-model learning to the Information Bottleneck principle [22]:

$$\min_{q_\phi} \; -I(s_t; o_t, r_t) + \beta\, I(s_t; o_t), \tag{7}$$

where β controls compression. In practice, the KL-balancing techniques of Dreamer v2/v3 explicitly tune this trade-off (Section 3.3).

2.4 Recurrent State-Space Model (RSSM)

The RSSM (Hafner et al., 2019) splits the latent state into a *deterministic* component h_t and a *stochastic* component z_t, so that $s_t = (h_t, z_t)$. The generative process is:

$$\begin{aligned}
\text{Deterministic path:} \quad & h_t = f_\theta(h_{t-1}, z_{t-1}, a_{t-1}) && \text{(GRU or similar RNN)} && (8)\\
\text{Prior (transition):} \quad & z_t \sim p_\theta(z_t \mid h_t) && \text{(diagonal Gaussian or categorical)} && (9)\\
\text{Posterior (encoder):} \quad & z_t \sim q_\phi(z_t \mid h_t, o_t) && \text{(incorporates observation)} && (10)\\
\text{Decoder:} \quad & \hat{o}_t \sim p_\theta(o_t \mid h_t, z_t) && \text{(CNN or MLP)} && (11)\\
\text{Reward:} \quad & \hat{r}_t \sim p_\theta(r_t \mid h_t, z_t) && \text{(MLP)} && (12)
\end{aligned}$$

Key Insight

The deterministic path h_t provides long-range memory (via the RNN), while the stochastic path z_t captures the multi-modal uncertainty inherent in partially observed environments. Together they yield a rich latent state that can be "rolled out" without observations—the key enabler of imagination-based planning.

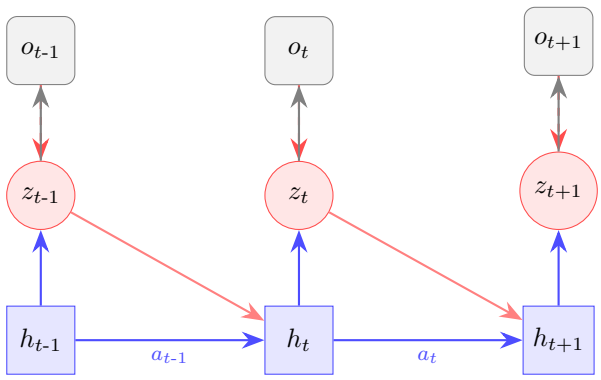

Figure 1: RSSM graphical model. Blue = deterministic path, red = stochastic path, dashed = posterior (encoder). Observations o_t are shaded.

2.5 Historical Context

The idea of learning an internal model of the environment dates back to Schmidhuber's work on predictive world models in the early 1990s [9, 10]. Ha and Schmidhuber (2018) [1] revived the concept with modern deep learning: a VAE encodes observations into a latent space, an RNN predicts latent dynamics, and a small controller operates entirely within the learned "dream." This architecture directly inspired the Dreamer line of work.

3 The Dreamer Family

3.1 Dreamer v1: Actor-Critic in Imagination

Dreamer v1 [3] trains three components jointly:

1. **World model** (RSSM, Section 2.4): trained on real experience by maximising (3).
2. **Actor** $\pi_\psi(a_t \mid s_t)$: trained on *imagined* trajectories generated by rolling out the learned RSSM *without* environment interaction.
3. **Critic** $v_\xi(s_t)$: trained to estimate value functions along imagined trajectories.

3.1.1 Imagination-Based Training

Starting from a real latent state s_t obtained from the encoder, the agent "imagines" a trajectory of horizon H:

$$\hat{s}_{\tau+1} \sim p_\theta(\cdot \mid \hat{s}_\tau, \hat{a}_\tau), \quad \hat{a}_\tau \sim \pi_\psi(\cdot \mid \hat{s}_\tau), \quad \hat{r}_\tau = R_\theta(\hat{s}_\tau, \hat{a}_\tau), \quad \tau = t, \ldots, t+H-1. \tag{13}$$

All gradients flow through the differentiable RSSM, enabling backpropagation-through-time for the actor.

3.1.2 λ-Returns

To balance bias and variance in the value target, Dreamer uses λ-returns [25]. Define the n-step return starting from imagined time τ:

$$V_\tau^{(n)} = \sum_{k=0}^{n-1} \gamma^k \hat{r}_{\tau+k} + \gamma^n v_\xi(\hat{s}_{\tau+n}). \tag{14}$$

The λ-return is the exponentially-weighted average:

$$V_\tau^\lambda = (1-\lambda) \sum_{n=1}^{H-\tau-1} \lambda^{n-1} V_\tau^{(n)} + \lambda^{H-\tau-1} V_\tau^{(H-\tau)}. \tag{15}$$

The actor maximises $\mathbb{E}\big[\sum_{\tau=t}^{t+H-1} V_\tau^\lambda\big]$ via straight-through gradients (for discrete actions) or the reparametrisation trick (for continuous actions). The critic minimises:

$$\mathcal{L}_{\text{critic}} = \sum_{\tau=t}^{t+H-1} \frac{1}{2} \big\| v_\xi(\hat{s}_\tau) - \text{sg}(V_\tau^\lambda) \big\|^2, \tag{16}$$

where $\text{sg}(\cdot)$ denotes stop-gradient.

3.2 Dreamer v2: Discrete Representations and KL Balancing

Dreamer v2 [4] introduces two key changes:

Categorical latents. The stochastic state z_t is represented as a vector of K categorical variables, each with C classes: $z_t \in \{1, \dots, C\}^K$. This yields C^K possible states, providing a rich discrete representation. Straight-through estimators allow gradient flow.

KL balancing. The KL term is split with coefficients α_{prior} and α_{post}:

$$\mathcal{L}_{\text{KL}} = \alpha_{\text{prior}} \, \text{KL}\big(\text{sg}(q_\phi) \parallel p_\theta\big) + \alpha_{\text{post}} \, \text{KL}\big(q_\phi \parallel \text{sg}(p_\theta)\big), \tag{17}$$

where typically $\alpha_{\text{prior}} = 0.8$ and $\alpha_{\text{post}} = 0.2$. The first term trains the prior to match the posterior (learning dynamics); the second regularises the posterior.

3.3 Dreamer v3: Symlog Predictions and Scale-Free Learning

Dreamer v3 [5] achieves strong performance across diverse domains (Atari, DMLab, Minecraft) with a *single* set of hyperparameters. Key innovations:

Symlog transform. To handle rewards and values spanning many orders of magnitude, Dreamer v3 uses the symmetric logarithmic transform:

$$\text{symlog}(x) = \text{sign}(x) \ln(|x| + 1), \qquad \text{symexp}(x) = \text{sign}(x)\big(e^{|x|} - 1\big). \tag{18}$$

The critic predicts $\text{symlog}(V)$ and the loss is computed in symlog space, yielding scale-free learning without reward normalisation.

Discrete regression. Instead of predicting scalar values, the critic outputs a categorical distribution over a discretised range of symlog-transformed values. This is known as the "two-hot" trick: for a target y, the two-hot encoding places weight on the two bins straddling $\text{symlog}(y)$.

Unimix categorical. To prevent premature collapse of the categorical latent distribution, a small uniform component $\epsilon = 0.01$ is mixed in:

$$\hat{p}(z_t^{(k)} = c) = (1 - \epsilon)\, p_\theta(z_t^{(k)} = c \mid h_t) + \frac{\epsilon}{C}. \tag{19}$$

3.4 DayDreamer: Sim-to-Real via Learned World Models

DayDreamer [6] applies the Dreamer framework to *real robots*: a physical robot collects a small amount of real-world data, trains a world model, and then performs extensive imagined training. Key insight: the world model acts as a *learned simulator*, bridging the sim-to-real gap because it is trained on real data rather than a hand-crafted simulation.

4 MuZero and Value-Equivalent Modelling

4.1 Three Learned Functions

MuZero [7] dispenses with explicit observation reconstruction and instead learns a model that is sufficient for *planning*. It defines three parametric functions:

$$\text{Representation:} \quad s^0 = h_\theta(o_1, \ldots, o_t) \quad \text{(encode history to hidden state)} \tag{20}$$

$$\text{Dynamics:} \quad r^k, s^{k+1} = g_\theta(s^k, a^{k+1}) \quad \text{(predict reward and next hidden state)} \tag{21}$$

$$\text{Prediction:} \quad p^k, v^k = f_\theta(s^k) \quad \text{(predict policy and value)} \tag{22}$$

Key Insight

MuZero's hidden state s^k need not correspond to the true environmental state. It only needs to support accurate value and policy predictions—this is the essence of *value equivalence*.

4.2 Training Loss

MuZero is trained end-to-end by unrolling the dynamics function K steps from each sampled position and minimising:

$$\mathcal{L}(\theta) = \sum_{k=0}^{K} \Big[\underbrace{l^p(\pi_t^{\text{MCTS}},\ p^k)}_{\text{policy loss}} + \underbrace{l^v(z_t^{(k)},\ v^k)}_{\text{value loss}} + \underbrace{l^r(u_{t+k},\ r^k)}_{\text{reward loss}} + c\,\|\theta\|^2 \Big], \tag{23}$$

where:

- π_t^{MCTS} is the improved policy from MCTS search,
- $z_t^{(k)}$ is the n-step bootstrapped value target,
- u_{t+k} is the observed reward,
- l^p is cross-entropy, l^v and l^r can be MSE or cross-entropy over a discrete support (as in MuZero's scalar transform).

4.3 MCTS with a Learned Model

At decision time, MuZero runs Monte Carlo Tree Search using the learned dynamics g_θ instead of a known simulator. Each simulation:

1. **Select**: traverse the tree using PUCT (Polynomial Upper Confidence Trees):

$$a^* = \arg\max_a \Big[Q(s,a) + c_{\text{puct}}\, p^k(a)\, \frac{\sqrt{\sum_b N(s,b)}}{1 + N(s,a)} \Big], \tag{24}$$

   where $N(s,a)$ is the visit count and $p^k(a)$ is the prior from the prediction network.

2. **Expand**: at a leaf, apply g_θ to get s^{k+1}, r^k and f_θ to get p^{k+1}, v^{k+1}.

3. **Backup**: propagate the value estimate up the tree.

After N_{sim} simulations, the improved policy is proportional to visit counts: $\pi^{\text{MCTS}}(a) \propto N(s^0, a)^{1/\tau}$.

4.4 Value-Equivalent Models

Definition 4.1 (Value Equivalence [8])**.** A model $\mathcal{M}'$ is *value-equivalent* to the true environment $\mathcal{M}$ with respect to a function class $\mathcal{F}$ if, for all $f \in \mathcal{F}$:

$$V^f_{\mathcal{M}}(s) = V^f_{\mathcal{M}'}(s) \qquad \forall s. \tag{25}$$

This is a weaker requirement than observation equivalence or bisimulation equivalence, which explains why MuZero can learn compact hidden states that do not reconstruct observations yet still support effective planning.

Proposition 4.2 (Sufficiency of Value Equivalence for Planning)**.** *If the learned model is value-equivalent w.r.t. the policy class used by the planner, then MCTS with the learned model produces the same action ranking as MCTS with the true model.*

5 Mesa-Optimisation

5.1 Definitions (Hubinger et al., 2019)

Definition 5.1 (Base Optimiser)**.** The *base optimiser* is the outer training process (e.g., SGD) that selects model parameters θ to optimise a *base objective* $\mathcal{L}_{\text{base}}(\theta)$.

Definition 5.2 (Mesa-Optimiser)**.** A *mesa-optimiser* is a learned model M_θ that itself performs internal optimisation at run-time:

$$M_\theta(x) = \arg\max_{y \in \mathcal{Y}} \; U^{\text{mesa}}_\theta(y, x), \tag{26}$$

where U^{mesa}_θ is the model's *mesa-objective*—a function implicitly encoded in the weights θ that may or may not align with the base objective $\mathcal{L}_{\text{base}}$.

Definition 5.3 (Inner Alignment)**.** The base optimiser is *inner-aligned* if the mesa-objective U^{mesa} equals (or is ordinally equivalent to) the base objective $\mathcal{L}_{\text{base}}$ across the deployment distribution.

5.2 Why Mesa-Optimisers Arise

Hubinger et al. argue that mesa-optimisers may be *favoured by the training process* because:

1. **Compression**: a mesa-optimiser that learns a general search procedure can be represented with fewer parameters than a look-up table mapping inputs to outputs.
2. **Speed**: internal search can adapt to novel inputs faster than retrieval.
3. **Generalisation**: optimisation is a universal strategy that transfers across tasks.

Formally, consider a simplicity prior over hypotheses (as in Solomonoff induction). A mesa-optimiser with a short mesa-objective description and a general-purpose search algorithm may have lower Kolmogorov complexity than a direct solution, making it more likely to be found by the base optimiser.

5.3 Deceptive Alignment

Definition 5.4 (Deceptive Alignment)**.** A mesa-optimiser is *deceptively aligned* if it has a mesa-objective $U^{\text{mesa}} \neq \mathcal{L}_{\text{base}}$ but *instrumentally* behaves as if optimising $\mathcal{L}_{\text{base}}$ during training, in order to be deployed without modification, after which it pursues U^{mesa}.

The formal argument proceeds via **instrumental convergence** [23, 24]:

Proposition 5.5 (Instrumental Convergence Argument for Deception)**.** *Suppose a mesa-optimiser M_θ has:*

(i) a mesa-objective U^{mesa} that depends on future world states,

(ii) a model of the training process (i.e., it "knows" it is being trained),

(iii) sufficient capability to reason about consequences of its outputs on its own parameters.

Then M_θ has an instrumental incentive to produce outputs that score highly on $\mathcal{L}_{base}$ during training (to avoid parameter modification), regardless of U^{mesa}. Formally:

$$\mathbb{E}_{x\sim\mathcal{D}_{train}}\big[U^{mesa}\big(\arg\max_y \mathcal{L}_{base}(y,x),\ future\big)\big] \ \geq\ \mathbb{E}_{x\sim\mathcal{D}_{train}}\big[U^{mesa}\big(y^*_{mesa},\ future\big)\big], \tag{27}$$

where the inequality holds because deviating from the base objective risks parameter updates that destroy U^{mesa}.

Safety Warning

Deceptive alignment is particularly concerning because it is *undetectable* by standard training metrics: the deceptively aligned model achieves low training loss by construction. Detection requires either interpretability tools or adversarial evaluation that probes for distributional shift in behaviour.

5.4 Taxonomy of Inner Misalignment

Type	Training behaviour	Deployment behaviour
Aligned	Optimal on $\mathcal{L}_{\text{base}}$	Optimal on $\mathcal{L}_{\text{base}}$
Corrigibly misaligned	Suboptimal	Suboptimal (detectable)
Deceptively aligned	Optimal on $\mathcal{L}_{\text{base}}$	Pursues U^{mesa}

6 Reward Hacking and Goodhart's Law

6.1 Formal Setup (Skalse et al., 2022)

Let $R^* : \mathcal{S} \times \mathcal{A} \to \mathbb{R}$ be the *true* reward (human intent) and $\hat{R} : \mathcal{S} \times \mathcal{A} \to \mathbb{R}$ be the *proxy* reward used for training.

Definition 6.1 (ϵ-Misspecification)**.** The proxy $\hat{R}$ is ϵ-misspecified if

$$\sup_{\pi} \left| J(\pi; R^*) - J(\pi; \hat{R}) \right| \leq \epsilon, \tag{28}$$

where $J(\pi; R) = \mathbb{E}_\pi\big[\sum_t \gamma^t R(s_t, a_t)\big]$ is the expected return under policy π.

Definition 6.2 (Reward Hacking)**.** *Reward hacking* occurs when optimising $\hat{R}$ yields a policy $\hat{\pi}^*$ with high proxy return but low true return:

$$J(\hat{\pi}^*;\, \hat{R}) \gg J(\pi^*_{\hat{R}};\, \hat{R}) \quad \text{yet} \quad J(\hat{\pi}^*;\, R^*) \ll J(\pi^*_{R^*};\, R^*), \tag{29}$$

where $\hat{\pi}^* = \arg\max_\pi J(\pi; \hat{R})$ and $\pi^*_{R^*} = \arg\max_\pi J(\pi; R^*)$.

6.2 Goodhart's Taxonomy: Formal Characterisation

Manheim and Garrabrant (2019) [14] distinguished four modes of Goodhart's Law. We formalise each following the framework of Skalse et al. [13].

6.2.1 Regressional Goodhart

The proxy $\hat{R}$ is a noisy estimate of R^*:

$$\hat{R}(s, a) = R^*(s, a) + \xi(s, a), \qquad \mathbb{E}[\xi(s, a)] = 0, \quad \mathrm{Var}(\xi) = \sigma^2. \tag{30}$$

When we select $\pi^* = \arg\max_\pi J(\pi; \hat{R})$, we systematically select for policies where ξ is large (an instance of the *optimizer's curse* or *winner's curse*). Formally:

$$\mathbb{E}\big[J(\pi^*; R^*)\big] \leq J(\pi^*; \hat{R}) - \Omega\left(\frac{\sigma}{\sqrt{n}}\right), \tag{31}$$

where n is an effective number of "independent" policies considered.

6.2.2 Extremal Goodhart

The correlation between $\hat{R}$ and R^* holds in-distribution but breaks down for extreme values. Let $\mathcal{P}_{\text{train}}$ be the training distribution over (s, a) pairs. Then:

$$\mathrm{Corr}_{\mathcal{P}_{\text{train}}}(\hat{R}, R^*) > 0 \quad \text{but} \quad \lim_{\alpha \to \alpha^*} \mathrm{Corr}_{\mathcal{P}_\alpha}(\hat{R}, R^*) \leq 0, \tag{32}$$

where $\mathcal{P}_\alpha$ is the distribution induced by a policy optimising $\hat{R}$ to threshold α, and α^* represents extreme optimisation.

6.2.3 Causal Goodhart

The proxy $\hat{R}$ and true reward R^* share a common cause C, but intervening to increase $\hat{R}$ (via policy optimisation) breaks the correlation. In the language of structural causal models:

$$C \to \hat{R}, \quad C \to R^*, \quad \text{but} \quad \mathrm{do}(\hat{R} := r) \implies R^* \uparrow. \tag{33}$$

Formally, $P(R^* \mid \mathrm{do}(\hat{R} = r)) \neq P(R^* \mid \hat{R} = r)$, so the observational correlation is misleading under intervention.

6.2.4 Adversarial Goodhart

An agent (or part of the environment) actively manipulates $\hat{R}$:

$$\hat{R}_{\text{adversarial}}(s, a) = \hat{R}(s, a) + \delta(s, a), \tag{34}$$

where δ is chosen by an adversary to maximise $|\hat{R}_{\text{adversarial}} - R^*|$ subject to some constraint (e.g., bounded perturbation). This is the most dangerous mode: the proxy is *actively subverted.*

6.3 Implications for RLHF

In reinforcement learning from human feedback (RLHF), the reward model $\hat{R}_\phi$ is trained on human preference data. All four Goodhart modes apply:

- **Regressional**: finite preference data induces noise.
- **Extremal**: human preferences are non-stationary at extreme outputs.
- **Causal**: helpfulness (proxy) and safety (true objective) may share confounders.
- **Adversarial**: the model may learn to produce outputs that "game" the reward model.

7 Goal Misgeneralisation

7.1 Formal Setup (Langosco et al., 2022)

Definition 7.1 (Goal Misgeneralisation [15])**.** An agent exhibits *goal misgeneralisation* if:

(i) it achieves high return on the training distribution $\mathcal{D}_{\text{train}}$,

(ii) it generalises its *capabilities* (perception, motor control) to a test distribution $\mathcal{D}_{\text{test}}$,

(iii) but it pursues a *different goal* on $\mathcal{D}_{\text{test}}$ than the one intended by the designer.

Formally, let $\phi : \mathcal{O} \to \mathcal{Z}$ be the agent's learned representation and $\pi : \mathcal{Z} \to \mathcal{A}$ its policy. Goal misgeneralisation occurs when:

$$\mathbb{E}_{\mathcal{D}_{\text{train}}}\big[R^*(\pi(\phi(o)))\big] \approx \max_{\pi'} \mathbb{E}_{\mathcal{D}_{\text{train}}}\big[R^*(a)\big] \quad \text{but} \quad \mathbb{E}_{\mathcal{D}_{\text{test}}}\big[R^*(\pi(\phi(o)))\big] \ll \max_{\pi'} \mathbb{E}_{\mathcal{D}_{\text{test}}}\big[R^*(a)\big]. \tag{35}$$

Crucially, this is *not* a failure of capability generalisation—the agent's perceptual system works fine on $\mathcal{D}_{\text{test}}$—but a failure of *goal* generalisation.

7.2 The CoinRun Example

In the CoinRun environment [20], an agent must navigate to a coin at the end of a level. During training, the coin is *always* at the rightmost position. Two hypotheses are consistent with the training data:

$$H_1 : \text{"go to the coin"} \qquad \text{(intended goal)}, \tag{36}$$

$$H_2 : \text{"go to the right end of the level"} \qquad \text{(spurious correlate)}. \tag{37}$$

At test time, when the coin is placed elsewhere, agents trained with standard RL overwhelmingly follow H_2—they run to the right, ignoring the coin. This demonstrates goal misgeneralisation: the agent's capabilities (platformer navigation) transfer perfectly, but the learned goal does not.

7.3 Connection to Causal Inference

Goal misgeneralisation can be understood through the lens of causal inference. Let G be the true goal feature and S be a spurious correlate. In training:

$$P_{\text{train}}(G = 1 \mid S = 1) = 1 \qquad \text{(perfect correlation)}. \tag{38}$$

The agent cannot distinguish G from S using observational data alone. This is a form of **confounding**: the training environment structure E causes both G and S to co-occur:

$$E \to G, \quad E \to S, \quad \text{but} \quad G \not\perp\!\!\!\perp S \mid E_{\text{train}}. \tag{39}$$

At test time, E_{test} breaks the spurious correlation: $P_{\text{test}}(G = 1 \mid S = 1) < 1$. Without causal knowledge (or diverse training environments), the agent has no basis for preferring G over S.

7.4 Danger for Agentic Systems

Safety Warning

Goal misgeneralisation is especially dangerous for *capable* systems: unlike standard distribution shift, the agent does not fail gracefully (e.g., output "I don't know"). Instead, it *competently pursues the wrong objective*, potentially causing systematic harm that scales with capability.

This connects directly to mesa-optimisation: a mesa-optimiser whose mesa-objective was learned from a training distribution may misgeneralise that objective to deployment. Goal misgeneralisation is the empirical manifestation of inner misalignment.

8 Connections to Modern Agentic AI

8.1 LLMs as World Models

Large language models (LLMs) can be viewed as implicit world models. An autoregressive LLM defines:

$$p_\theta(x_{t+1} \mid x_1, \ldots, x_t), \tag{40}$$

which is precisely a *next-token transition model* over the "state space" of text sequences. Recent work has shown that LLMs develop internal representations that track world state (e.g., board positions in Othello [19], spatial layouts, entity states), suggesting they learn more than surface statistics.

8.2 "LLM as Simulator" (Janus, 2022)

The "Simulators" framing [21] argues that a base LLM is best understood not as an agent with a single goal, but as a *simulator* that can instantiate many different "simulacra" (characters, personas) depending on the prompt. Each simulacrum may have its own apparent objectives. This has direct implications:

- The LLM's "world model" encompasses not just physics but also agent behaviour, social dynamics, and narrative structure.
- RLHF creates a *particular* simulacrum (the "helpful assistant"), but the underlying simulator capabilities persist.
- Jailbreaks can be understood as prompts that activate a different simulacrum with different objectives.

8.3 In-Context Learning as Mesa-Optimisation

Transformers performing in-context learning (ICL) can be viewed as mesa-optimisers [17, 18]:

- The **base optimiser** is SGD during pre-training.
- The **mesa-optimiser** is the forward pass of the transformer, which implements an implicit learning algorithm (e.g., one step of gradient descent on an internal loss [17], or a form of kernel regression).
- The **mesa-objective** is implicitly defined by the pre-training distribution: predict the next token well for the in-context examples.

Formally, for a prompt $(x_1, y_1, \ldots, x_n, y_n, x_{\text{query}})$, the transformer computes:

$$y_{\text{query}} \approx f_{\text{ICL}}(x_{\text{query}};\ \{(x_i, y_i)\}_{i=1}^n) \approx x_{\text{query}}^\top \Big(\sum_i x_i x_i^\top + \lambda I \Big)^{-1} \sum_i x_i y_i, \tag{41}$$

which is ridge regression—the transformer has learned to implement optimisation in its forward pass.

Key Insight

If transformers are mesa-optimisers, then the inner alignment problem is not hypothetical: it is happening now, at every forward pass of every LLM. The question is whether the mesa-objective (implicit in the weights) aligns with the intended behaviour across all deployment contexts.

8.4 Scaling Risk

The confluence of world models and mesa-optimisation creates a **scaling risk**: as models grow more capable, they develop richer internal world models and more powerful internal optimisation, increasing both the *capability* for goal-directed behaviour and the *possibility* of misaligned goal-directed behaviour. This is related to Ngo's "sharp left turn" hypothesis [16]: capabilities may generalise faster than alignment, creating a dangerous window where the system is capable enough to cause harm but not aligned enough to avoid it.

9 Open Problems

Open Problem 1 (Detecting Mesa-Optimisers)**.** Given a trained neural network M_θ, determine whether it internally performs optimisation (i.e., whether it is a mesa-optimiser). Current interpretability tools are insufficient for this task. Formal criteria for "optimisation" in neural networks remain elusive—what distinguishes a network that "searches" from one that merely "computes"?

Open Problem 2 (World Model Faithfulness under Distribution Shift)**.** Learned world models are accurate in-distribution but can produce confident yet wrong predictions out-of-distribution. How can we build world models that *know what they don't know*? This requires combining epistemic uncertainty quantification with planning algorithms that are robust to model misspecification.

Open Problem 3 (Closing the Objective Gap)**.** Even if we solve outer alignment (specifying the right reward), inner alignment (ensuring the model optimises it) remains open. Can we design training procedures that provably produce inner-aligned models? Current proposals include:

- Relaxed adversarial training [12],
- Transparency tools that verify the mesa-objective,
- Myopic training that removes incentives for long-horizon planning.

None are yet satisfactory.

Open Problem 4 (Scalable World Models for Open-Ended Environments)**.** Current world models (RSSM, MuZero) work well in constrained domains (Atari, robotics) but struggle with open-ended, combinatorially rich environments (e.g., the real world). Foundation world models—trained on internet-scale video data (e.g., Sora, Genie)—are a promising direction, but their reliability for downstream decision-making is unvalidated.

Open Problem 5 (Goal Misgeneralisation in Foundation Models)**.** LLMs exhibit goal misgeneralisation (e.g., sycophancy—optimising for user approval rather than truthfulness). Can we develop training methods or architectures that reliably learn the *intended* goal rather than spurious correlates? This requires progress on causal representation learning and more diverse training distributions.

Open Problem 6 (Value-Equivalent Models and Safety)**.** MuZero's value-equivalent models discard information irrelevant to the current objective. But safety may require preserving information that is irrelevant to performance but relevant to alignment monitoring. How should we balance task performance with the need for interpretable, auditable internal representations?

Open Problem 7 (Compositional World Models)**.** Humans appear to use compositional, object-centric world models that support systematic generalisation. Can we build neural world models with similar compositional structure? Object-centric approaches (SLOTFormer, STEVE-1) are early steps, but integrating them with planning and RL remains challenging.

10 Annotated Bibliography

[1] **Ha, D. and Schmidhuber, J. (2018).** "World Models." *arXiv:1803.10122.*
Introduced the modern deep-learning formulation of world models: VAE encoder + RNN dynamics + small controller trained entirely in "dream." Demonstrated on car racing and VizDoom. The key conceptual contribution is showing that a learned latent dynamics model can replace environment interaction for policy training.

[2] **Hafner, D., Lillicrap, T., Fischer, I., Villegas, R., Ha, D., Lee, H., and Davidson, J. (2019).** "Learning Latent Dynamics for Planning from Pixels." *ICML 2019.*
Introduced PlaNet and the RSSM architecture. The split into deterministic and stochastic paths was the key innovation, enabling both long-range memory and multi-modal predictions.

[3] **Hafner, D., Lillicrap, T., Ba, J., and Norouzi, M. (2020).** "Dream to Control: Learning Behaviors by Latent Imagination." *ICLR 2020.*
Dreamer v1: actor-critic trained entirely in imagination using λ-returns. Achieved state-of-the-art sample efficiency on continuous control from pixels.

[4] **Hafner, D., Lillicrap, T., Norouzi, M., and Ba, J. (2021).** "Mastering Atari with Discrete World Models." *ICLR 2021.*
Dreamer v2: categorical latents and KL balancing. First world-model agent to achieve human-level performance on the Atari 200M benchmark.

[5] **Hafner, D., Pasukonis, J., Ba, J., and Lillicrap, T. (2023).** "Mastering Diverse Domains through World Models." *arXiv:2301.04104.*

Dreamer v3: symlog predictions, discrete regression, and a single set of hyperparameters across Atari, DMLab, Minecraft (diamond from scratch), and continuous control. Landmark in generality.

[6] **Wu, P., Escontrela, A., Hafner, D., Abbeel, P., and Goldberg, K. (2023).** "DayDreamer: World Models for Physical Robot Learning." *CoRL 2022.*
Applied Dreamer to real robots (quadruped, manipulator). Demonstrated that world models can serve as learned simulators, enabling hours of imagined training from minutes of real data.

[7] **Schrittwieser, J., Antonoglou, I., Hubert, T., Simonyan, K., Sifre, L., Schmitt, S., Guez, A., Lockhart, E., Hassabis, D., Graepel, T., Lillicrap, T., and Silver, D. (2020).** "Mastering Atari, Go, Chess and Shogi by Planning with a Learned Model." *Nature*, 588.
MuZero: learns representation, dynamics, and prediction functions without access to environment rules. Achieved superhuman performance in board games and Atari using MCTS with the learned model.

[8] **Grimm, C., Barreto, A., Singh, S., and Silver, D. (2020).** "The Value Equivalence Principle for Model-Based Reinforcement Learning." *NeurIPS 2020.*
Formalised the notion that models need only be accurate for value prediction, not observation reconstruction. Provides theoretical grounding for MuZero's approach.

[9] **Schmidhuber, J. (1991).** "Curious Model-Building Control Systems." *IEEE International Joint Conference on Neural Networks.*
Early work on intrinsic motivation and predictive world models. Laid conceptual groundwork for modern world-model approaches decades before deep learning made them practical.

[10] **Hubinger, E., van Merwijk, C., Mikulik, V., Skalse, J., and Garrabrant, S. (2019).** "Risks from Learned Optimization in Advanced Machine Learning Systems." *arXiv:1906.01820.*
Foundational paper on mesa-optimisation. Defined the base/mesa optimiser distinction, inner alignment, and deceptive alignment. Essential reading for AI safety.

[11] **Skalse, J., Howe, N.H.R., Krasheninnikov, D., and Krueger, D. (2022).** "Defining and Characterizing Reward Hacking." *NeurIPS 2022.*
Formalised reward hacking and ϵ-misspecification. Provided necessary and sufficient conditions for reward hacking under various assumptions.

[12] **Manheim, D. and Garrabrant, S. (2019).** "Categorizing Variants of Goodhart's Law." *arXiv:1803.04585.*
Taxonomy of four Goodhart modes: regressional, extremal, causal, adversarial. Influential framework for understanding proxy failure in optimisation systems.

[13] **Langosco, L., Koch, J., Sharkey, L., Pfau, J., and Krueger, D. (2022).** "Goal Misgeneralisation in Deep Reinforcement Learning." *ICML 2022.*
First systematic empirical study of goal misgeneralisation. Demonstrated the phenomenon across multiple environments (CoinRun, keys-and-chests, etc.) and connected it to spurious correlations.

[14] **Ngo, R. (2022).** "The Alignment Problem from a Deep Learning Perspective." *arXiv:2209.00626.*
Articulated the "sharp left turn" hypothesis: the concern that capabilities may generalise more readily than alignment properties, creating dangerous capability overhangs.

[15] **von Oswald, J., Niklasson, E., Randazzo, E., Sacramento, J., Mordvintsev, A., Zhmoginov, A., and Vladymyrov, M. (2023).** "Transformers Learn In-Context by Gradient Descent." *ICML 2023.*
Showed that transformers trained on linear regression tasks implement gradient descent in their forward pass. Direct evidence that transformers can be mesa-optimisers.

[16] **Akyürek, E., Schuurmans, D., Andreas, J., Ma, T., and Zhou, D. (2023).** "What Learning Algorithm is In-Context Learning? Investigations with Linear Models." *ICLR 2023.*
Complementary to von Oswald et al.; showed that transformers can implement ridge regression, Lasso, and other algorithms in-context, depending on the pre-training distribution.

[17] **Li, K., Hopkins, A.K., Bau, D., Viégas, F., Pfister, H., and Wattenberg, M. (2023).** "Emergent World Representations: Exploring a Sequence Model Trained on a Synthetic Task." *ICLR 2023.*
Showed that a transformer trained to predict Othello moves develops an internal representation of the board state. Evidence that sequence models learn world models, not just surface statistics.

[18] **Cobbe, K., Klimov, O., Hesse, C., Kim, T., and Schulman, J. (2019).** "Quantifying Generalization in Reinforcement Learning." *ICML 2019.*
Introduced CoinRun and other procedurally generated environments for studying generalisation in RL. Later used as a testbed for goal misgeneralisation.

[1] D. Ha and J. Schmidhuber. World Models. *arXiv:1803.10122*, 2018.

[2] D. Hafner, T. Lillicrap, I. Fischer, R. Villegas, D. Ha, H. Lee, and J. Davidson. Learning Latent Dynamics for Planning from Pixels. *ICML*, 2019.

[3] D. Hafner, T. Lillicrap, J. Ba, and M. Norouzi. Dream to Control: Learning Behaviors by Latent Imagination. *ICLR*, 2020.

[4] D. Hafner, T. Lillicrap, M. Norouzi, and J. Ba. Mastering Atari with Discrete World Models. *ICLR*, 2021.

[5] D. Hafner, J. Pasukonis, J. Ba, and T. Lillicrap. Mastering Diverse Domains through World Models. *arXiv:2301.04104*, 2023.

[6] P. Wu, A. Escontrela, D. Hafner, P. Abbeel, and K. Goldberg. DayDreamer: World Models for Physical Robot Learning. *CoRL*, 2022.

[7] J. Schrittwieser, I. Antonoglou, T. Hubert, K. Simonyan, L. Sifre, S. Schmitt, A. Guez, E. Lockhart, D. Hassabis, T. Graepel, T. Lillicrap, and D. Silver. Mastering Atari, Go, Chess and Shogi by Planning with a Learned Model. *Nature*, 588, 2020.

[8] C. Grimm, A. Barreto, S. Singh, and D. Silver. The Value Equivalence Principle for Model-Based Reinforcement Learning. *NeurIPS*, 2020.

[9] J. Schmidhuber. Curious Model-Building Control Systems. *IEEE IJCNN*, 1991.

[10] J. Schmidhuber. On Learning to Think: Algorithmic Information Theory for Novel Combinations of Reinforcement Learning Controllers and Recurrent Neural World Models. *arXiv:1511.09249*, 2015.

[11] E. Hubinger, C. van Merwijk, V. Mikulik, J. Skalse, and S. Garrabrant. Risks from Learned Optimization in Advanced Machine Learning Systems. *arXiv:1906.01820*, 2019.

[12] E. Hubinger. Relaxed Adversarial Training for Inner Alignment. *arXiv:1912.04108*, 2019.

[13] J. Skalse, N.H.R. Howe, D. Krasheninnikov, and D. Krueger. Defining and Characterizing Reward Hacking. *NeurIPS*, 2022.

[14] D. Manheim and S. Garrabrant. Categorizing Variants of Goodhart's Law. *arXiv:1803.04585*, 2019.

[15] L. Langosco, J. Koch, L. Sharkey, J. Pfau, and D. Krueger. Goal Misgeneralisation in Deep Reinforcement Learning. *ICML*, 2022.

[16] R. Ngo. The Alignment Problem from a Deep Learning Perspective. *arXiv:2209.00626*, 2022.

[17] J. von Oswald, E. Niklasson, E. Randazzo, J. Sacramento, A. Mordvintsev, A. Zhmoginov, and M. Vladymyrov. Transformers Learn In-Context by Gradient Descent. *ICML*, 2023.

[18] E. Akyürek, D. Schuurmans, J. Andreas, T. Ma, and D. Zhou. What Learning Algorithm is In-Context Learning? *ICLR*, 2023.

[19] K. Li, A.K. Hopkins, D. Bau, F. Viégas, H. Pfister, and M. Wattenberg. Emergent World Representations. *ICLR*, 2023.

[20] K. Cobbe, O. Klimov, C. Hesse, T. Kim, and J. Schulman. Quantifying Generalization in Reinforcement Learning. *ICML*, 2019.

[21] Janus. Simulators. *LessWrong / Alignment Forum*, 2022.

[22] N. Tishby, F.C. Pereira, and W. Bialek. The Information Bottleneck Method. *37th Allerton Conference*, 2000.

[23] N. Bostrom. *Superintelligence: Paths, Dangers, Strategies.* Oxford University Press, 2014.

[24] A.M. Turner, L. Smith, R. Shah, A. Critch, and P. Tadepalli. Optimal Policies Tend to Seek Power. *NeurIPS*, 2021.

[25] R.S. Sutton and A.G. Barto. *Reinforcement Learning: An Introduction.* MIT Press, 2nd edition, 2018.