Thousands of AI Authors

What 2,778 researchers actually think the next 20 years look like

after Grace, Stewart, Sandkühler, Thomas, Weinstein-Raun, Brauner (2024)

One survey. One question per expert.

2,778

AI researchers

every NeurIPS, ICML, ICLR, JMLR, AAAI, IJCAI author from the previous year

Asked in 2023.

Same questions as 2022.

Largest expert survey of its kind. Answers sealed before GPT-4 retail launch.

The story of this paper is four moments.

Each one makes you stop.

1. A thirteen-year shift.
2. A dot plot that already aged badly.
3. A tail that will not go away.
4. What they actually fear.

Moment 1 — The thirteen-year shift

The question

When will unaided machines accomplish every task better and more cheaply than human workers?

They call it **High-Level Machine Intelligence.**
Call it AGI if you prefer.

In 2022, the aggregate answer was:

2060

50% probability date for HLMI

Roughly a generation away.

One year later, the same experts said:

2047

50% probability date for HLMI

Thirteen years closer. In twelve months.

wait —

WHAT?

An expert consensus shifted by 13 years. In one year.

The shift, as a picture

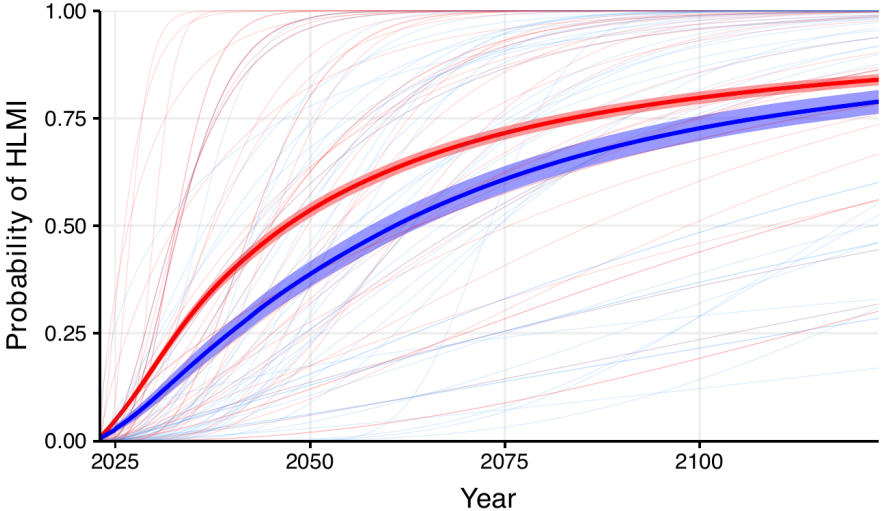

2023 Aggregate Forecast (with 95% Confidence Interval)

And for “all human jobs”

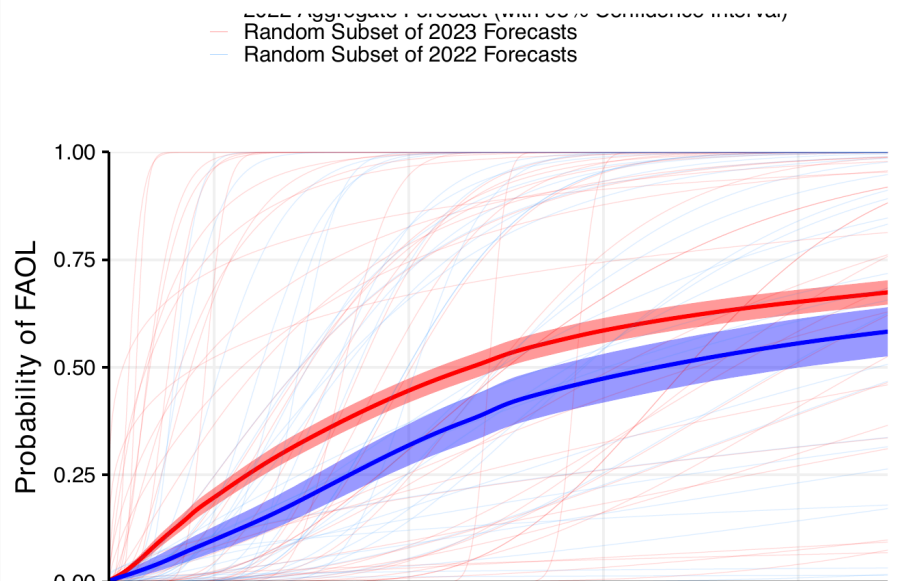

Read that again

The 50% date for automating
every job on earth
moved from 2164 to 2116.

That is not a rounding error.

Moment 2 — A dot plot that already aged badly

“HLMI” is too abstract.

So they asked about 39 specific things.

Write a *New York Times* best-seller.
Fold laundry.
Translate into a newfound language.
Do your taxes. Play Starcraft from screen video.

For each: when does 50% probability cross?

Task
Job
2022
2023
Year
2050
2100
2150
Full Automation of Labor (all human jobs)
AI Researcher
High Level Machine Intelligence (all human tasks)
Surgeon
Publishable math theorems
NYT best-selling fiction
Translate text in newfound language
Retail Salesperson
Win Putnam Math Competition
Simple Python code given spec and examples
Answers undecided questions well
Answer open-ended fact questions with web
Explain AI actions in games
5km city race as bipedal robot (superhuman)
Equations governing virtual worlds
Top 40 Pop Song
Translate text (vs. fluent amateur)
Transcribe speech
Answer factoid questions with web
Beat humans at Go (after same # games)
Recognize object seen once
Fake new song by specific artist
Translate speech using subtitles
Telephone banking services
Read text aloud
Random new computer game (novice level)
Construct video from new angle
Group new objects into classes
Good high school history essay
World Series of Poker
Atari games after 20m play (50% vs. novice)
Angry Birds (superhuman)
Learn to sort long lists
All Atari games (vs. pro game tester)
Fold laundry
Truck Driver
Assemble LEGO given instructions
Top Starcraft play via video of screen
Milestones
2025
2030
2035
Year

Look at the top of that chart

- **Write a NYT best-seller.** Forecast: 2033.
- **Answer open-ended factual questions.** Forecast: 2027.
- **Write simple Python from a spec.** Forecast: 2025.

Several of these were arguably done before the ink dried.

wait —

WHAT?

Expert forecasts for “hard” AI milestones now read like a to-do list.

Half the pace, inside a decade

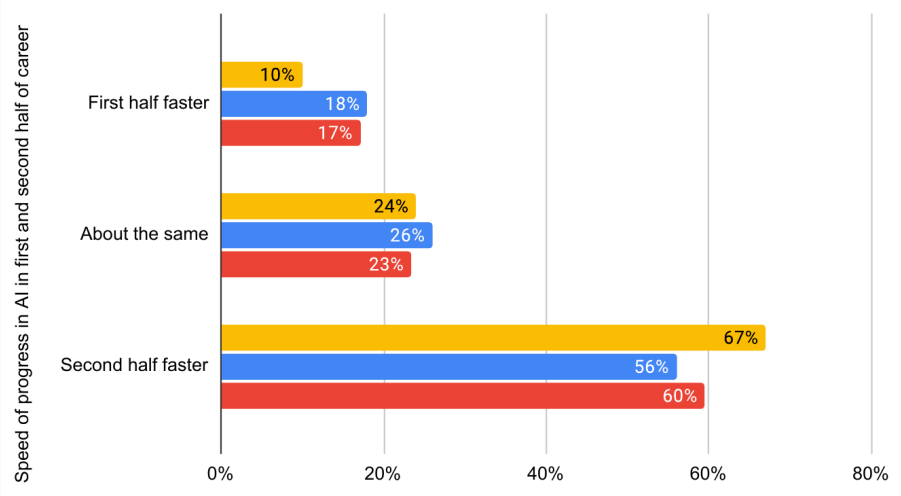

Aggregate probability each narrow task is feasible within N years.

The manager's question

Which of these dots
sits inside **your** core workflow?

That is the real due-diligence exercise.

Moment 3 — A tail that will not go away

They were also asked about outcomes.

"If HLMI arrives, how does it go?"

Five buckets, from *extremely good* to *extremely bad*.

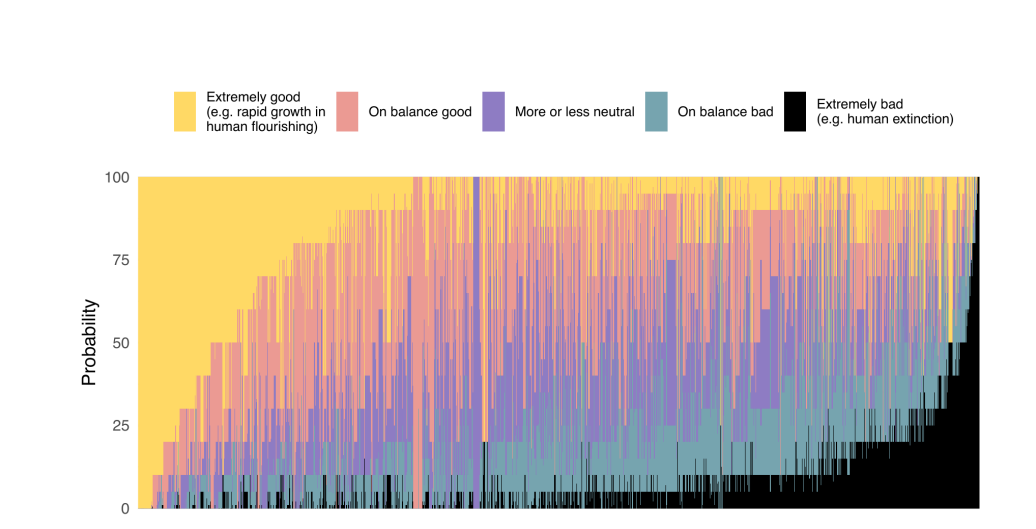

Each vertical bar is one respondent, sorted by optimism.

The number that stops a boardroom

38%

of surveyed AI researchers
give at least 10% probability
to outcomes *"as bad as human extinction"*.

wait —

WHAT?

Not a fringe. Not a tabloid. A plurality of the field.

The extinction-tail distribution

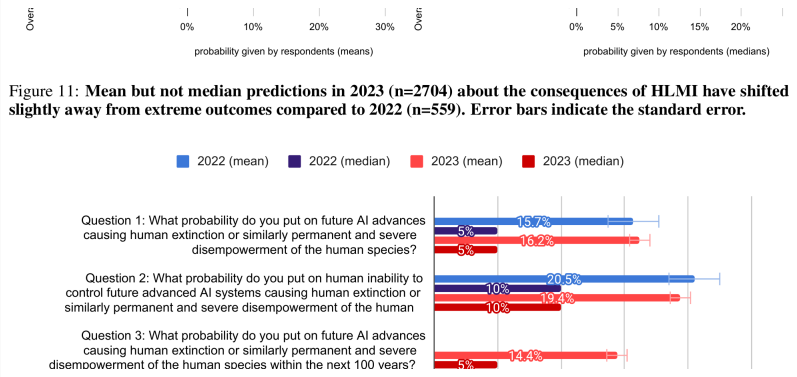

Figure 11: **Mean but not median predictions in 2023 (n=2704) about the consequences of HLMI have shifted slightly away from extreme outcomes compared to 2022 (n=559). Error bars indicate the standard error.**

Respondent-level probability assigned to extinction-grade outcomes.

And “intelligence explosion”

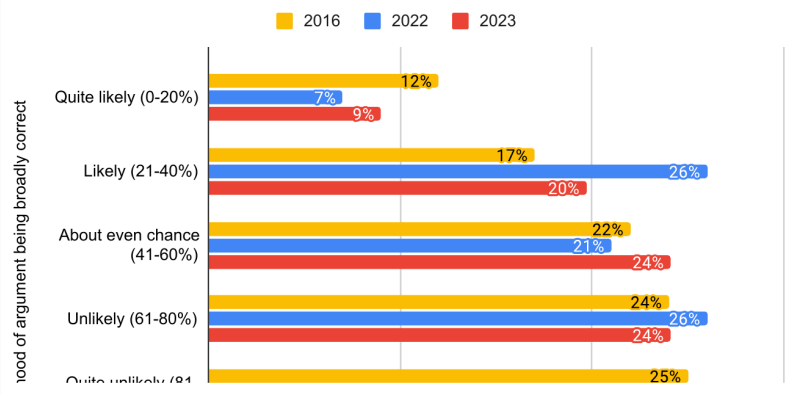

Likelihood that AI improves AI in a rapid, self-reinforcing loop.

The calibration question you should ask

"If a bridge engineer told you there's a 10% chance the bridge **collapses** — would you cross it?"

That is the shape of this tail.

Moment 4 — What they actually fear

Not robots. Not sentience.

Two behaviours keep coming up.

Deception.
Power-seeking.

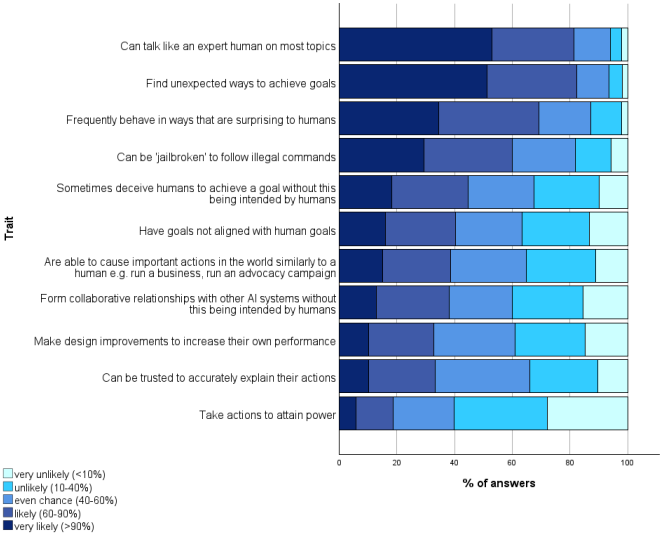
Trait
Can talk like an expert human on most topics
Find unexpected ways to achieve goals
Frequently behave in ways that are surprising to humans
Can be 'jailbroken' to follow illegal commands
Sometimes deceive humans to achieve a goal without this being intended by humans
Have goals not aligned with human goals
Are able to cause important actions in the world similarly to a human e.g. run a business, run an advocacy campaign
Form collaborative relationships with other AI systems without this being intended by humans
Make design improvements to increase their own performance
Can be trusted to accurately explain their actions
Take actions to attain power
0
20
40
60
80
100
% of answers
very unlikely (<10%)
unlikely (10-40%)
even chance (40-60%)
likely (60-90%)
very likely (>90%)

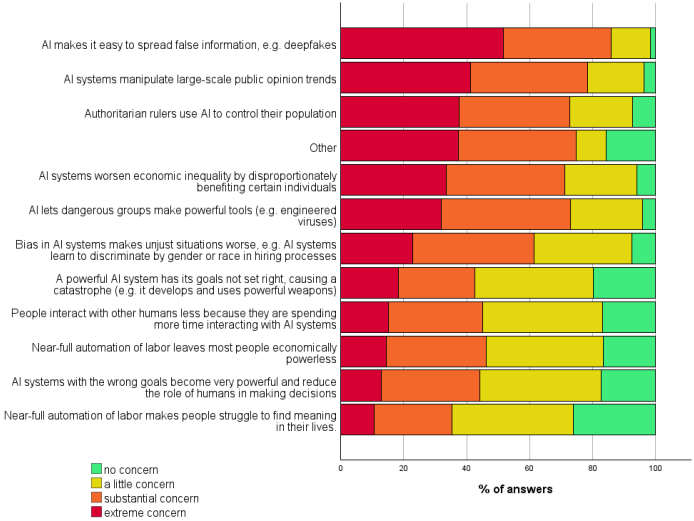
AI makes it easy to spread false information, e.g. deepfakes
AI systems manipulate large-scale public opinion trends
Authoritarian rulers use AI to control their population
Other
AI systems worsen economic inequality by disproportionately benefiting certain individuals
AI lets dangerous groups make powerful tools (e.g. engineered viruses)
Bias in AI systems makes unjust situations worse, e.g. AI systems learn to discriminate by gender or race in hiring processes
A powerful AI system has its goals not set right, causing a catastrophe (e.g. it develops and uses powerful weapons)
People interact with other humans less because they are spending more time interacting with AI systems
Near-full automation of labor leaves most people economically powerless
AI systems with the wrong goals become very powerful and reduce the role of humans in making decisions
Near-full automation of labor makes people struggle to find meaning in their lives.
0
20
40
60
80
100
% of answers
no concern
a little concern
substantial concern
extreme concern

wait —

WHAT?

The concern is not science fiction. It is mis-optimisation at scale.

And what they want spent on it

3. Frequently behave in ways that are surprising to humans (69.1% of respondents)

Will AI in 2028 truthfully and intelligibly explain its decisions?

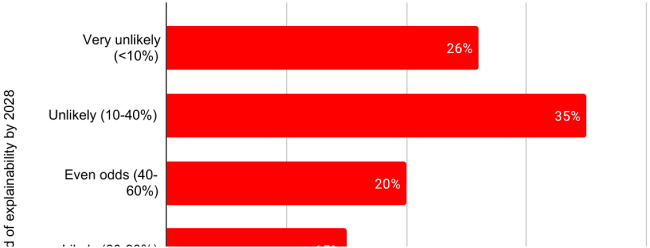

Majorities want more investment in AI safety / alignment research.

So what?

Three things to leave this room with

1. Expert timelines are **unstable**.
   Do not anchor a five-year plan on a point estimate.

2. Narrow-task forecasts **already under-shot**.
   Assume your workflow-level dot arrives early, not late.

3. The tail is **not a fringe**.
   38% of the field assigns $\geq$10% to catastrophic outcomes.

The one number

13

years of expected progress
compressed into *twelve months* of calendar time.

Plan as if that could happen again.

Thank you.

Questions.

Grace, K. *et al.* (2024).
Thousands of AI Authors on the Future of AI.
arXiv:2401.02843.